## **Module V**

# Process-to-Process Delivery: UDP, TCP

### 23-1 PROCESS-TO-PROCESS DELIVERY

The transport layer is responsible for process-toprocess delivery—the delivery of a packet, part of a message, from one process to another. Two processes communicate in a client/server relationship, as we will see later.

Topics discussed in this

**section:** Client/Server Paradigm

**Multiplexing and Demultiplexing** 

**Connectionless Versus Connection-Oriented** 

Service Reliable Versus Unreliable

**Three Protocols** 



The transport layer is responsible for process-to-process delivery.

### Figure 23.1 Types of data

#### deliveries

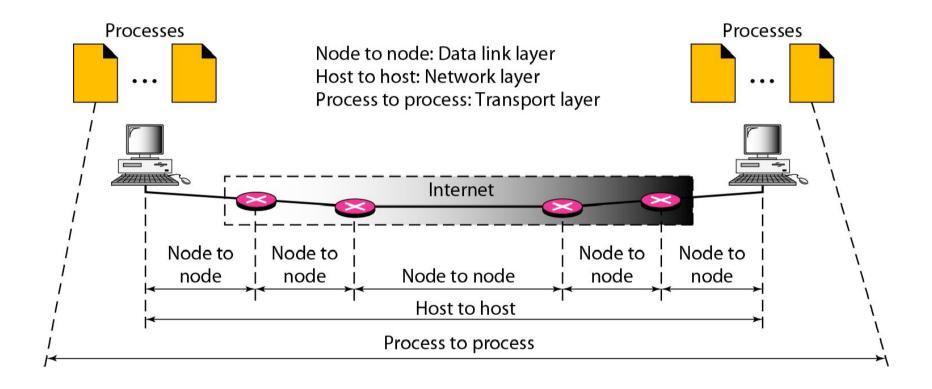


Figure 23.2 Port

numbers

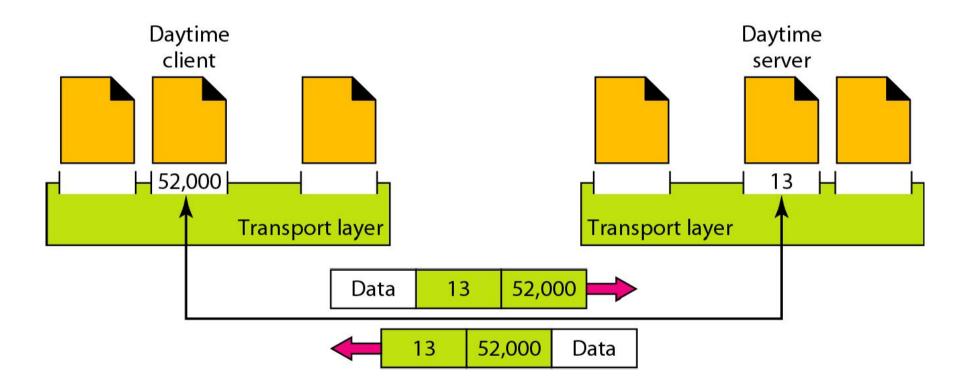
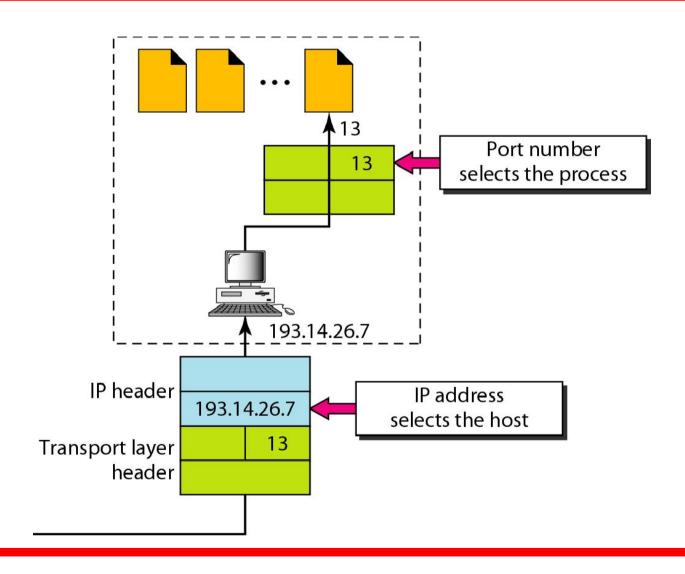


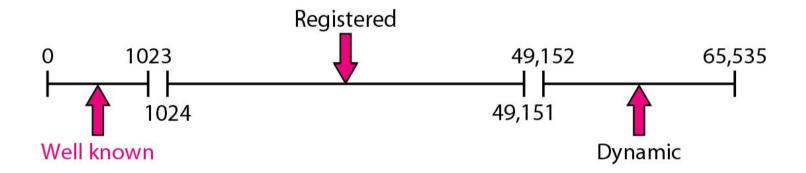
Figure 23.3 IP addresses versus port

numbers



### Figure 23.4 IANA

ranges



### Figure 23.5 Socket

address

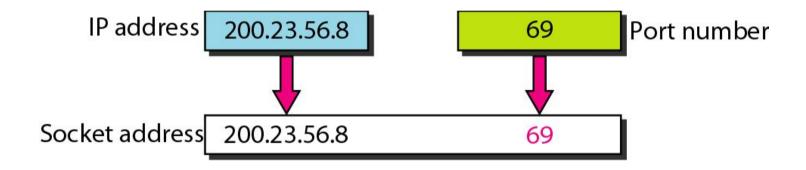
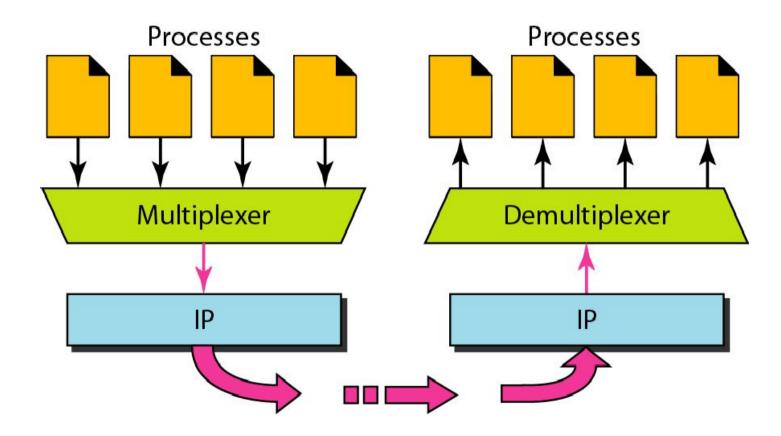


Figure 23.6 Multiplexing and demultiplexing



### Figure 23.7 Error

control

Error is checked in these paths by the data link layer

Error is not checked in these paths by the data link layer

Transport

Network

Data link

Physical

WAN

LAN

LAN

LAN

LAN

Figure 23.8 Position of UDP, TCP, and SCTP in TCP/IP

suite **Application SMTP** DNS **SNMP BOOTP** FTP **TFTP** layer Transport **TCP SCTP** UDP layer **IGMP** ICMP Network IΡ layer **ARP RARP** Data link layer **Underlying LAN or WAN** technology Physical layer

### 23-2 USER DATAGRAM PROTOCOL (UDP)

The User Datagram Protocol (UDP) is called a connectionless, unreliable transport protocol. It does not add anything to the services of IP except to provide process-to-process communication instead of host-to-host communication.

### Topics discussed in this section:

Well-Known Ports for UDP User Datagram Checksum UDP Operation Use of UDP

### Table 23.1 Well-known ports used with UDP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
53	Nameserver	Domain Name Service
67	BOOTPs	Server port to download bootstrap information
68	BOOTPc	Client port to download bootstrap information
69	TFTP	Trivial File Transfer Protocol
111	RPC	Remote Procedure Call
123	NTP	Network Time Protocol
161	SNMP	Simple Network Management Protocol
162	SNMP	Simple Network Management Protocol (trap)



### Example 23.1

In UNIX, the well-known ports are stored in a file called /etc/services Each line in this file gives the name of the server and the well-known port number. We can use the grep utility to extract the line corresponding to the desired application. The following shows the port for FTP. Note that FTP can use port 21 with either UDP or TCP

\$ grep	ftp	/etc/services
ftp	21/	tcp
ftp	21/	udp



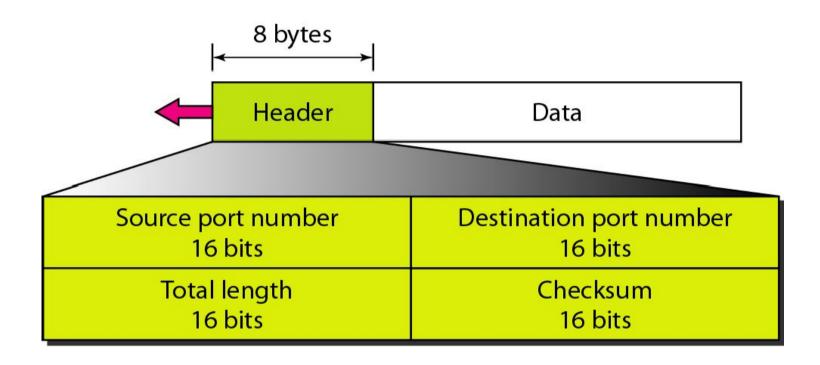
## Example 23.1 (continued)

SNMP uses two port numbers (161 and 162), each for a different purpose, as we will see in Chapter 28.

\$ grep	snmp /etc/services	
snmp	161/tcp	#Simple Net Mgmt Proto
snmp	161/udp	#Simple Net Mgmt Proto
snmptrap	162/udp	#Traps for SNMP

Figure 23.9 User datagram

format



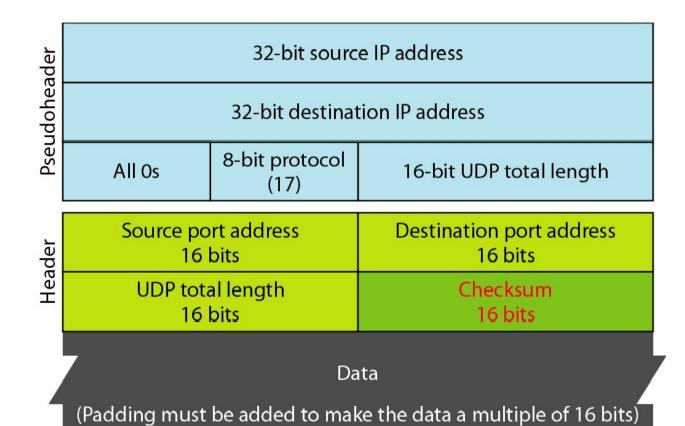


## **UDP** length

= IP length – IP header's length

### Figure 23.10 Pseudoheader for checksum

#### calculation



# Example 23.2

Figure 23.11 shows the checksum calculation for a very small user datagram with only 7 bytes of data. Because the number of bytes of data is odd, padding is added for checksum calculation. The pseudoheader as well as the padding will be dropped when the user datagram is delivered to IP.

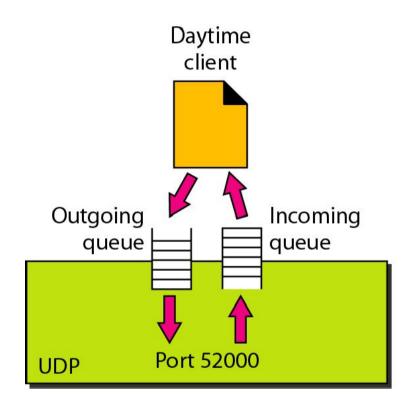
#### Figure 23.11 Checksum calculation of a simple UDP user

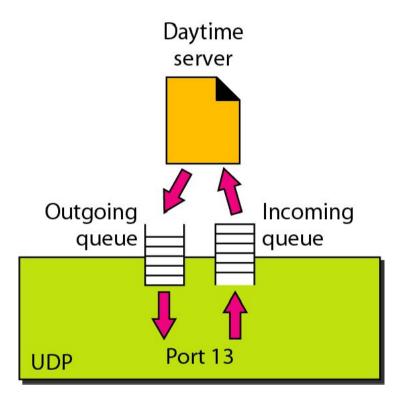
datagram

153.18.8.105					
171.2.14.10					
All Os	All Os 17 15				
10	87	13			
15		All Os			
Т	T E		Т		
L	N	G	All Os		

```
10011001 00010010 --- 153.18
00001000 01101001 --- 8.105
10101011 00000010 --- 171.2
00001110 00001010 --- 14.10
00000000 00010001 — → 0 and 17
00000100 00111111 --- 1087
00000000 00001101 --- 13
00000000 00000000 → 0 (checksum)
01010100 01000101 → Tand E
01010011 01010100 → Sand T
01001001 01001110 → land N
01000111 00000000 — F G and 0 (padding)
10010110 11101011 → Sum
01101001 00010100 ---- Checksum
```

Figure 23.12 Queues in UDP





### 23-3 TCP

TCP is a connection-oriented protocol; it creates a virtual connection between two TCPs to send data. In addition, TCP uses flow and error control mechanisms at the transport level.

### Topics discussed in this section:

**TCP Services** 

**TCP** 

**Features** 

**Segment** 

**A TCP** 

**Connection Flow** 

23.22 ontrol Error


**Control** 

### Table 23.2 Well-known ports used by TCP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
20	FTP, Data	File Transfer Protocol (data connection)
21	FTP, Control	File Transfer Protocol (control connection)
23	TELNET	Terminal Network
25	SMTP	Simple Mail Transfer Protocol
53	DNS	Domain Name Server
67	ВООТР	Bootstrap Protocol
79	Finger	Finger
80	HTTP	Hypertext Transfer Protocol
111	RPC	Remote Procedure Call

### Figure 23.13 Stream

delivery




### Figure 23.14 Sending and receiving

buffers Sending Receiving process process TCP **TCP** Next byte Next byte to write to read Buffer Buffer Not read Sent Not sent Next byte Next byte Stream of bytes to send to receive

### **Figure 23.15** *TCP*

segments





The bytes of data being transferred in each connection are numbered by TCP. The numbering starts with a randomly generated number.

# Example 23.3

The following showsthesequencenumber for each segment:

```
      Segment 1
      →
      Sequence Number: 10,001 (range: 10,001 to 11,000)

      Segment 2
      →
      Sequence Number: 11,001 (range: 11,001 to 12,000)

      Segment 3
      →
      Sequence Number: 12,001 (range: 12,001 to 13,000)

      Segment 4
      →
      Sequence Number: 13,001 (range: 13,001 to 14,000)

      Segment 5
      →
      Sequence Number: 14,001 (range: 14,001 to 15,000)
```



The value in the sequence number field of a segment defines the number of the first data byte contained in that segment.



The value of the acknowledgment field in a segment defines the number of the next byte a party expects to receive.

The acknowledgment number is cumulative.

### Figure 23.16 TCP segment

format Header Data Source port address Destination port address 16 bits 16 bits Sequence number 32 bits Acknowledgment number 32 bits U S **HLEN** Reserved Window size S 16 bits 4 bits 6 bits N Checksum Urgent pointer 16 bits 16 bits **Options and Padding** 

### Figure 23.17 Control

field

URG: Urgent pointer is valid

ACK: Acknowledgment is valid

PSH: Request for push

RST: Reset the connection

SYN: Synchronize sequence numbers

FIN: Terminate the connection

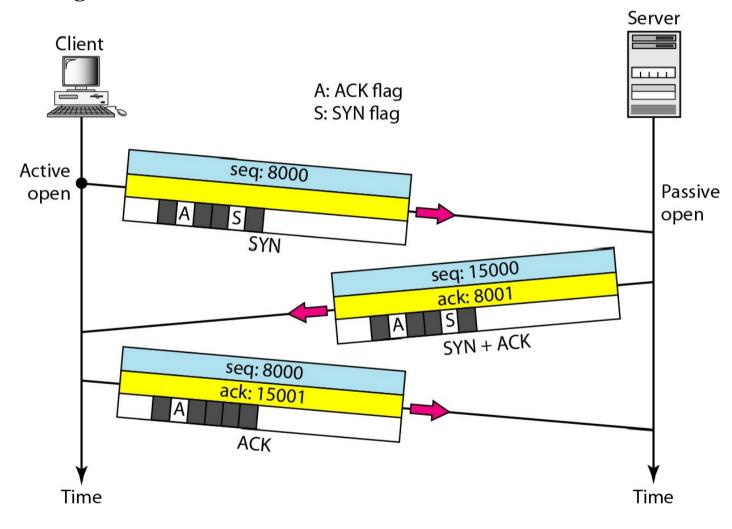
URG ACK PSH	RST	SYN	FIN
-------------	-----	-----	-----

Table 23.3Description of flags in the control field

Flag	Description
URG	The value of the urgent pointer field is valid.
ACK	The value of the acknowledgment field is valid.
PSH	Push the data.
RST	Reset the connection.
SYN	Synchronize sequence numbers during connection.
FIN	Terminate the connection.

Figure 23.18 Connection establishment using three-way

handshaking





A SYN segment cannot carry data, but it consumes one sequence number.



A SYN + ACK segment cannot carry data, but does consume one sequence number.

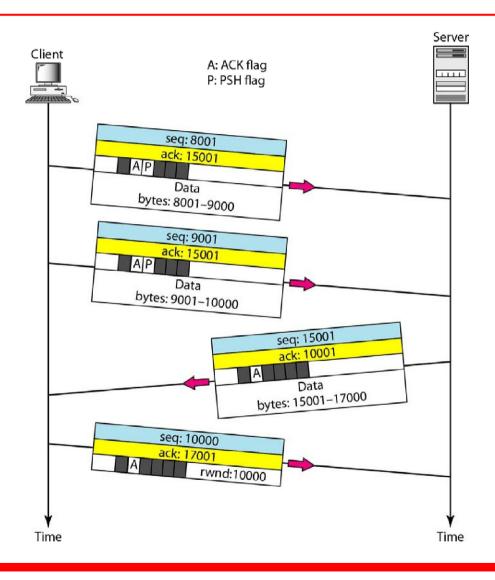


# An ACK segment, if carrying no data, consumes no sequence

number.

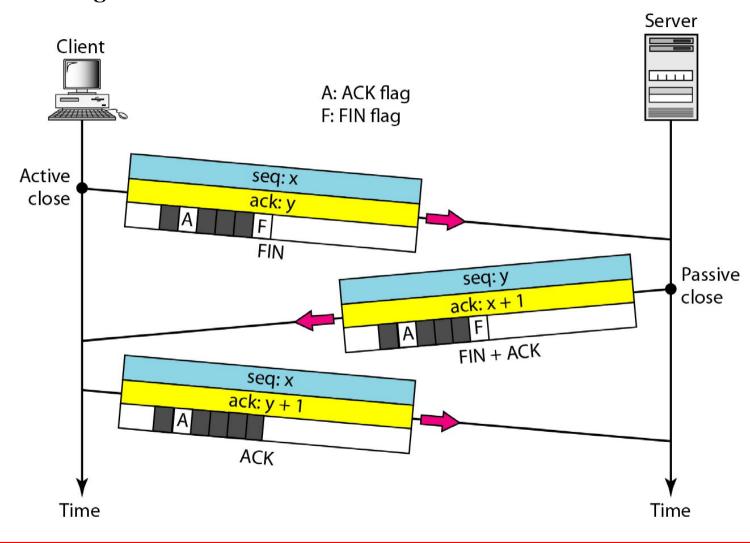
#### Figure 23.19 Data

transfer



#### Figure 23.20 Connection termination using three-way

handshaking





# The FIN segment consumes one sequence number if it does not carry data.

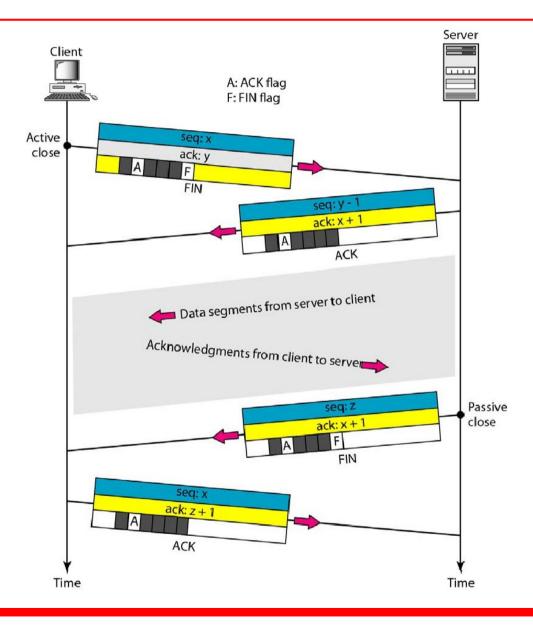


# The FIN + ACK segment consumes one sequence number if it

does not carry data.

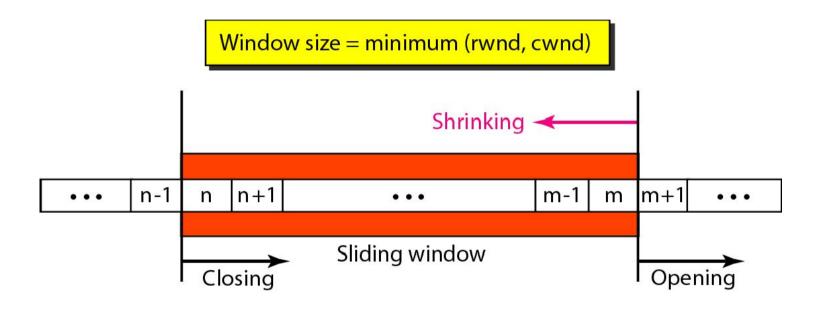
### **Figure 23.21**

#### Half-close



#### Figure 23.22 Sliding

window





A sliding window is used to make transmission more efficient as well as to control the flow of data so that the destination does not become overwhelmed with data.

TCP sliding windows are byte-oriented.

# Example 23.4

What is the value of the receiver window (rwnd) for host A if the receiver, host B, has a buffer size of 5000 bytes and 1000 bytes of received and unprocessed data?

#### Solution

The value of rwnd = 5000 - 1000 = 4000. Host B can receive only 4000 bytes of data before overflowing its buffer. Host B advertises this value in its next segment to A.



#### Example 23.5

What is the size of the window for host A if the value of rwnd is 3000 bytes and the value of cwnd is 3500 bytes?

#### Solution

The size of the window is the smaller of rwnd and cwnd, which is 3000 bytes.

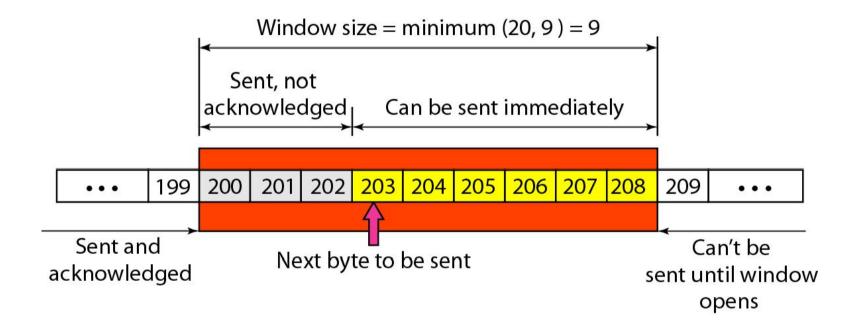


#### Example 23.6

Figure 23.23 shows an unrealistic example of a sliding window. The sender has sent bytes up to 202. We assume that cwnd is 20 (in reality this value is thousands of bytes). The receiver has sent an acknowledgment number of 200 with an rwnd of 9 bytes (in reality this value is thousands of bytes). The size of the sender window is the minimum of rwnd and cwnd, or 9 bytes. Bytes 200 to 202 are sent, but not acknowledged. Bytes 203 to 208 can be sent without worrying about acknowledgment. Bytes 209 and above cannot be sent.

Figure 23.23 Example

23.6





- The size of the window is the lesser of rwnd and cwnd.
- The source does not have to send a full window's worth of data.
- The window can be opened or closed by the receiver, but should not be shrunk.
- The destination can send an acknowledgment at any time as long as it does not result in a shrinking window.
- The receiver can temporarily shut down the window; the cender, however, can always cend a segment of 1 byte after the window is shut down.



# ACK segments do not consume sequence numbers and are not acknowledged.



In modern implementations, a retransmission occurs if the retransmission timer expires or three duplicate ACK segments have

arrivea.



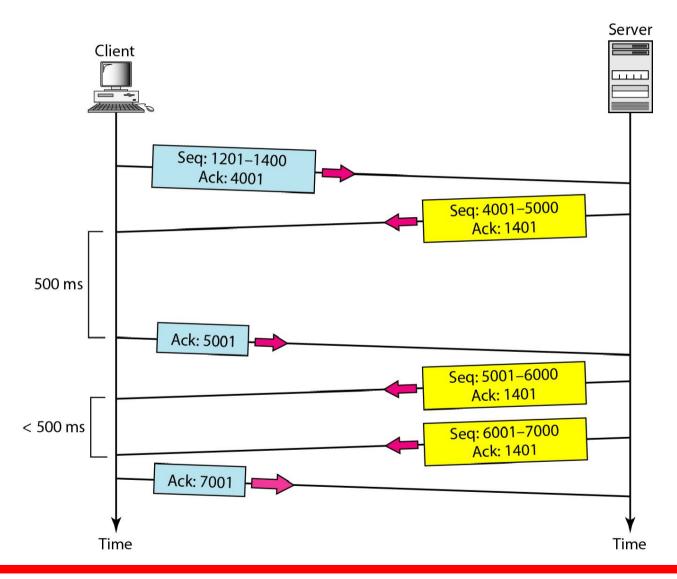
# No retransmission timer is set for an ACK segment.



Data may arrive out of order and be temporarily stored by the receiving TCP, but TCP guarantees that no out-of-order segment is delivered to the process.

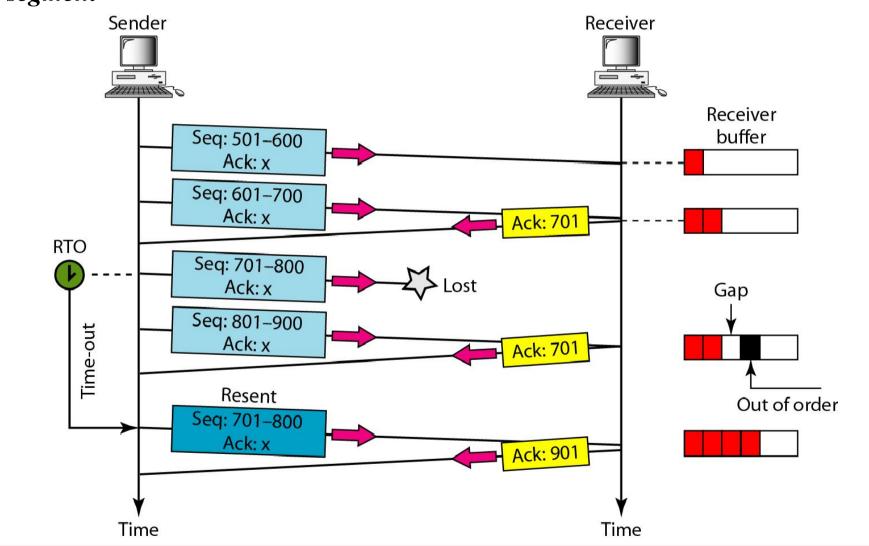
#### Figure 23.24 Normal

operation



#### Figure 23.25 Lost

segment





# The receiver TCP delivers only ordered data to the process.

#### Figure 23.26 Fast

#### retransmission

